

ELECTRA

ABSTRACT

INTRODUCTION

Learning Bidirectional Representation BERT
BERT Pre-training MLM NSP 가
MLM (15%) [MASK] (Noise Token) Encoder ,
Classifier (Linear Layer) Noise Token (Denoising;)
Encoder Denoising AutoEncoder .
가
(Representation) ,
Representation Learning .
BERT MLM task 15% BERT
Substantial Computing Cost()가 .

가
15% (Original Token + Noise Token)
Noise Token Loss .
Replaced Token Detection
MLM 가
Computationally Efficient() 가

가
Replaced Token Detection Masked Token Synthetically Generated Token(SGT
Token)
(Real;Not Replaced) 가 (Fake;Replaced)
SGT Token Generator MLM
MLM

We call our approach ELECTRA for “Efficiently Learning an Encoder that Classifies Token Replacements Accurately.”

(Replaced Token Detection)

BERT SOTA

•
•

METHOD

Replaced Token Detection Task

Generator	[MASK]	
Discriminator	SGT가	Real/Fake

Generator	Discriminator	Maximum Likelihood (MLE)	
가	가	KL-Div	

Cross Entropy KL-Div
One-Hot.(i.e All or None) , Cross Entropy 가 .(

Generator Loss

Generator	Representation Vector
-----------	-----------------------

$$p_G(x_t|\mathbf{x}) = \exp(e(x_t)^T h_G(\mathbf{x})_t) / \sum_{x'} \exp(e(x')^T h_G(\mathbf{x})_t)$$

x, x'	input token	.
h_G()	Generator	.
e()	Embedding	.

Generator, Look-up Table, Embedding Matrix, Embedding Vector, Generator Representation Vector, Header Network, Embedding Matrix.

$$\begin{aligned} h_G(x) &, \text{ vocab_size} \times \text{len}(\text{embedding_vector}) \\ e(x) &, 1 \times \text{len}(\text{embedding_vector}) \end{aligned}$$

$e(x) \cdot h_G(x)^T$, $1 \times \text{vocab_size}$
 Embedding Vector Embedding Matrix (i.e. Look-Up Table) ,
 . (, .)

Softmax

$$\mathcal{L}_{\text{MLM}}(\mathbf{x}, \theta_G) = \mathbb{E} \left(\sum_{i \in \mathbf{m}} -\log p_G(x_i | \mathbf{x}^{\text{masked}}) \right)$$

KL-Div 가 , One-Hot KL-Div Cross Entropy

Discriminator Loss

Discriminator	Fake/Real
---------------	-----------

$$D(\mathbf{x}, t) = \text{sigmoid}(w^T h_D(\mathbf{x})_t)$$

$$\mathcal{L}_{\text{Disc}}(\mathbf{x}, \theta_D) = \mathbb{E} \left(\sum_{t=1}^n -\mathbb{1}(x_t^{\text{corrupt}} = x_t) \log D(\mathbf{x}^{\text{corrupt}}, t) - \mathbb{1}(x_t^{\text{corrupt}} \neq x_t) \log(1 - D(\mathbf{x}^{\text{corrupt}}, t)) \right)$$

Discriminator (Binary)	Generator	Softmax	CrossEntropy
---------------------------	-----------	---------	--------------

We minimize the combined loss

$$\min_{\theta_G, \theta_D} \sum_{\mathbf{x} \in \mathcal{X}} \mathcal{L}_{\text{MLM}}(\mathbf{x}, \theta_G) + \lambda \mathcal{L}_{\text{Disc}}(\mathbf{x}, \theta_D)$$

lambda	Generator 가	Discriminator Discriminator	Loss 가	
Generator	Task()가	Discriminator	Task()	Entropy
Entropy		가	가	가
	BERT	Vocab_size entropy	30000 가 가	30000
	Discriminator	Weight	Hyper	Parameter

Method ,

- Replaced Token Detection Task
- Maximun Likelihood Estimation
- Discriminator Loss 가

Experiment

Experimental Setup

가 ()

General Language Understanding Evaluation(GLUE) Stanford Question Answering (SQuAD) dataset

GLUE

[[<https://huffon.github.io/2019/11/16/glue/>]]

SQuAD

[[<https://happygrammer.github.io/nlp/dataset/>]]

BERT

GLUE

Simple Linear Classifier Header Network

SQuAD

XL-NET Header Network

BERT Header Network(Independently predict) Jointly predict

(Model Extentions)

ELECTRA

BERT-base

가

가

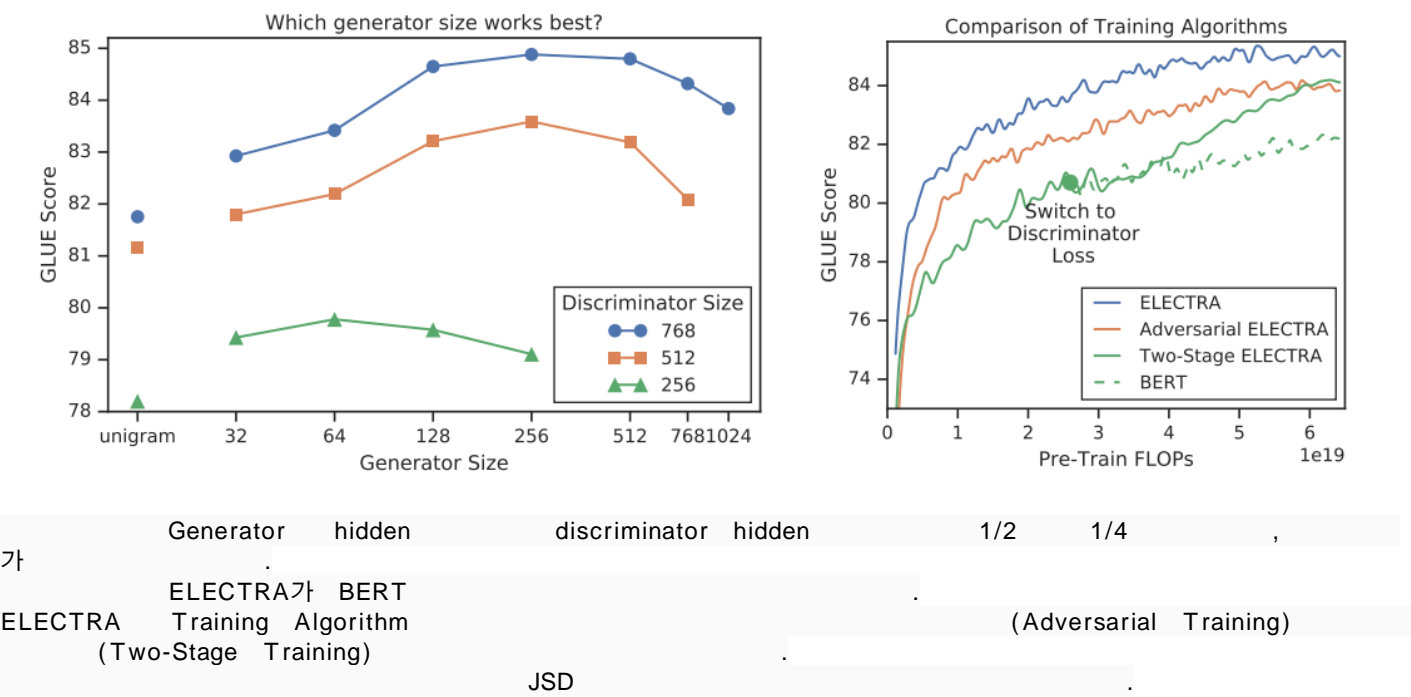
- Weight Sharing
 - Discriminator Generator Embedding Layer 가 (i.e Embedding Layer Look-Up table)
- Smaller Generator
 - ELECTRA, Generator Discriminator BERT
 - BERT, Generator Layer 가
 - Generator Unigram (Character Level)
 - 500k step
- Training Algorithm(Two-Step Training)
 - Generator Discriminator Generator n-step (, Generator 가)
 - Discriminator 가 n-step (, Generator 가)
 -)

SMALL MODEL LARGE MODEL

SMALL MODEL

As a goal of this work is to improve the efficiency of pre-training, we develop a small model that can be quickly trained on a single GPU.

SMALL



Model	Train / Infer FLOPs	Speedup	Params	Train Time + Hardware	GLUE
ELMo	3.3e18 / 2.6e10	19x / 1.2x	96M	14d on 3 GTX 1080 GPUs	71.2
GPT	4.0e19 / 3.0e10	1.6x / 0.97x	117M	25d on 8 P6000 GPUs	78.8
BERT-Small	1.4e18 / 3.7e9	45x / 8x	14M	4d on 1 V100 GPU	75.1
BERT-Base	6.4e19 / 2.9e10	1x / 1x	110M	4d on 16 TPuv3s	82.2
ELECTRA-Small	1.4e18 / 3.7e9	45x / 8x	14M	4d on 1 V100 GPU	79.9
50% trained	7.1e17 / 3.7e9	90x / 8x	14M	2d on 1 V100 GPU	79.0
25% trained	3.6e17 / 3.7e9	181x / 8x	14M	1d on 1 V100 GPU	77.7
12.5% trained	1.8e17 / 3.7e9	361x / 8x	14M	12h on 1 V100 GPU	76.0
6.25% trained	8.9e16 / 3.7e9	722x / 8x	14M	6h on 1 V100 GPU	74.1
ELECTRA-Base	6.4e19 / 2.9e10	1x / 1x	110M	4d on 16 TPuv3s	85.1

- 가

1. ELMo GPT , BERT-small , ELECTRA-small 가

2. BERT BERT

3. BERT BERT-small/base BERT

4. 1 GPU 1 GPU 가

LARGE MODEL

We train big ELECTRA models to measure the effectiveness of the replaced token detection pretraining task at the large scale of current state-of-the-art pre-trained Transformers.

LARGE Replaced Token Detection Task가 SOTA

Model	Train FLOPs	Params	SQuAD 1.1 dev		SQuAD 2.0 dev		SQuAD 2.0 test	
			EM	F1	EM	F1	EM	F1
BERT-Base	6.4e19 (0.09x)	110M	80.8	88.5	–	–	–	–
BERT	1.9e20 (0.27x)	335M	84.1	90.9	79.0	81.8	80.0	83.0
SpanBERT	7.1e20 (1x)	335M	88.8	94.6	85.7	88.7	85.7	88.7
XLNet-Base	6.6e19 (0.09x)	117M	81.3	–	78.5	–	–	–
XLNet	3.9e21 (5.4x)	360M	89.7	95.1	87.9	90.6	87.9	90.7
RoBERTa-100K	6.4e20 (0.90x)	356M	–	94.0	–	87.7	–	–
RoBERTa-500K	3.2e21 (4.5x)	356M	88.9	94.6	86.5	89.4	86.8	89.8
ALBERT	3.1e22 (44x)	235M	89.3	94.8	87.4	90.2	88.1	90.9
BERT (ours)	7.1e20 (1x)	335M	88.0	93.7	84.7	87.5	–	–
ELECTRA-Base	6.4e19 (0.09x)	110M	84.5	90.8	80.5	83.3	–	–
ELECTRA-400K	7.1e20 (1x)	335M	88.7	94.2	86.9	89.6	–	–
ELECTRA-1.75M	3.1e21 (4.4x)	335M	89.7	94.9	88.0	90.6	88.7	91.4

RoBERTa-500k	ELECTRA-400k	RoBERTa	ELECTRA	ELECTRA가 1/4	R
oBERTa					
ELECTRA-1.75M	RoBERTa-500k			가	
ELECTRA					

EFFICIENCY ANALYSIS

ELECTRA 가

가

- 1. Loss
- 2. Masked Language Model Task vs Replaced Token Detection Task
- 3. [MASK]

Model	ELECTRA	All-Tokens MLM	Replace MLM	ELECTRA 15%	BERT
GLUE score	85.0	84.3	82.4	82.4	82.2

Table 5: Compute-efficiency experiments (see text for details).

- 1. Loss
 - 2. [MASK]
 - 3. Replaced Token Detection Task Masked Language Model Task
- , ELECTRA 가

Screen Shot 2022-05-01 at 11.49.40 PM.png	55.4 KB	2022-05-01
Screen Shot 2022-05-02 at 12.41.59 AM.png	15 KB	2022-05-01
Screen Shot 2022-05-02 at 12.42.36 AM.png	14.9 KB	2022-05-01
Screen Shot 2022-05-02 at 12.42.36 AM.png	14.9 KB	2022-05-01
Screen Shot 2022-05-02 at 12.56.44 AM.png	14.9 KB	2022-05-01
Screen Shot 2022-05-02 at 1.01.26 AM.png	9.81 KB	2022-05-01
Screen Shot 2022-05-02 at 1.05.57 AM.png	20.2 KB	2022-05-01
Screen Shot 2022-05-02 at 1.11.42 AM.png	19.9 KB	2022-05-01
Screen Shot 2022-05-02 at 1.43.02 AM.png	106 KB	2022-05-01
Screen Shot 2022-05-02 at 1.49.15 AM.png	107 KB	2022-05-01
Screen Shot 2022-05-02 at 2.04.31 AM.png	115 KB	2022-05-01
Screen Shot 2022-05-02 at 2.38.35 AM.png	30.6 KB	2022-05-01